

Design

PSE of

PSE Group

Fraunhofer Institute for Systems and Innovation Research ISI
Advisor: M.Sc. Ankush Meshram

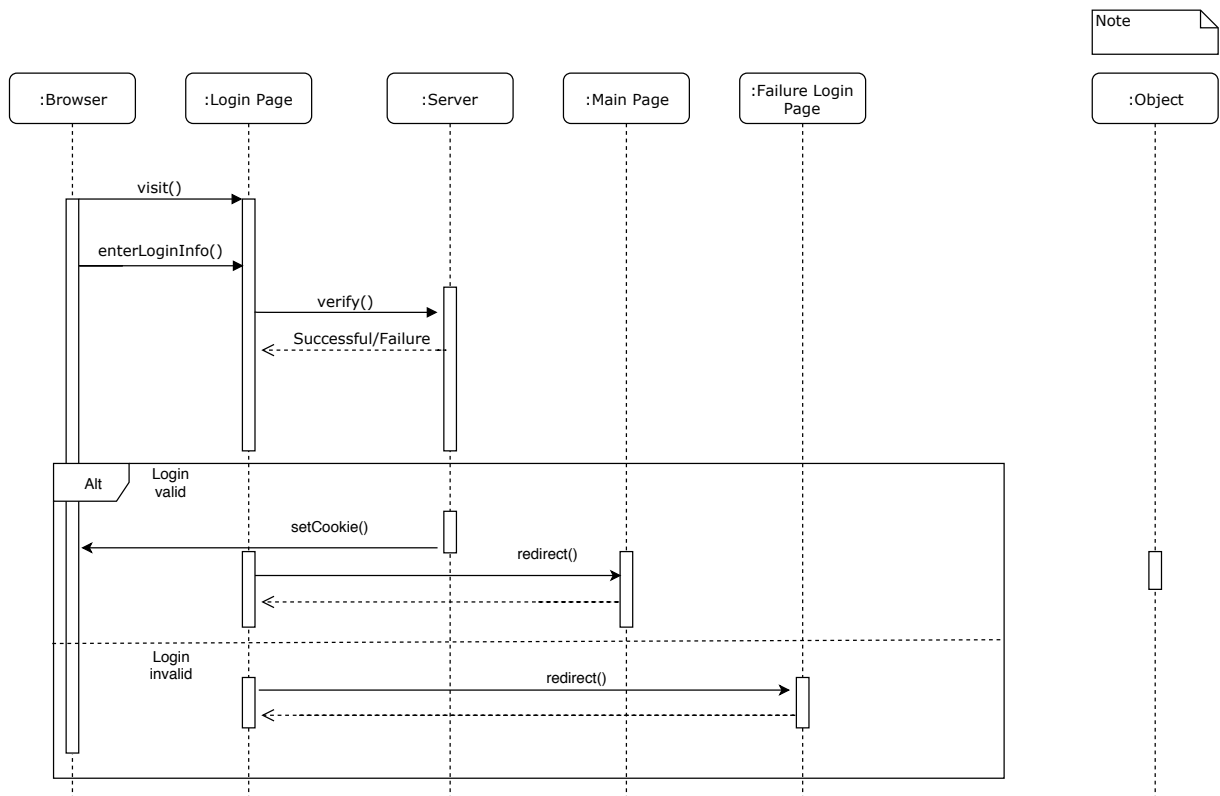
Contents

1	Design	1
1.1	Front-End	1
1.1.1	Sequence Diagram	1
1.1.2	Activity Diagram	5
1.1.3	Class Diagram	5
1.2	Client-server protocol	7
1.2.1	Requests from client to server:	7
1.2.2	Messages from server to client:	7
1.3	Back-End	9
1.3.1	Class Diagram	9
1.3.2	Sequence Diagram	20
1.3.3	Activity Diagram	20

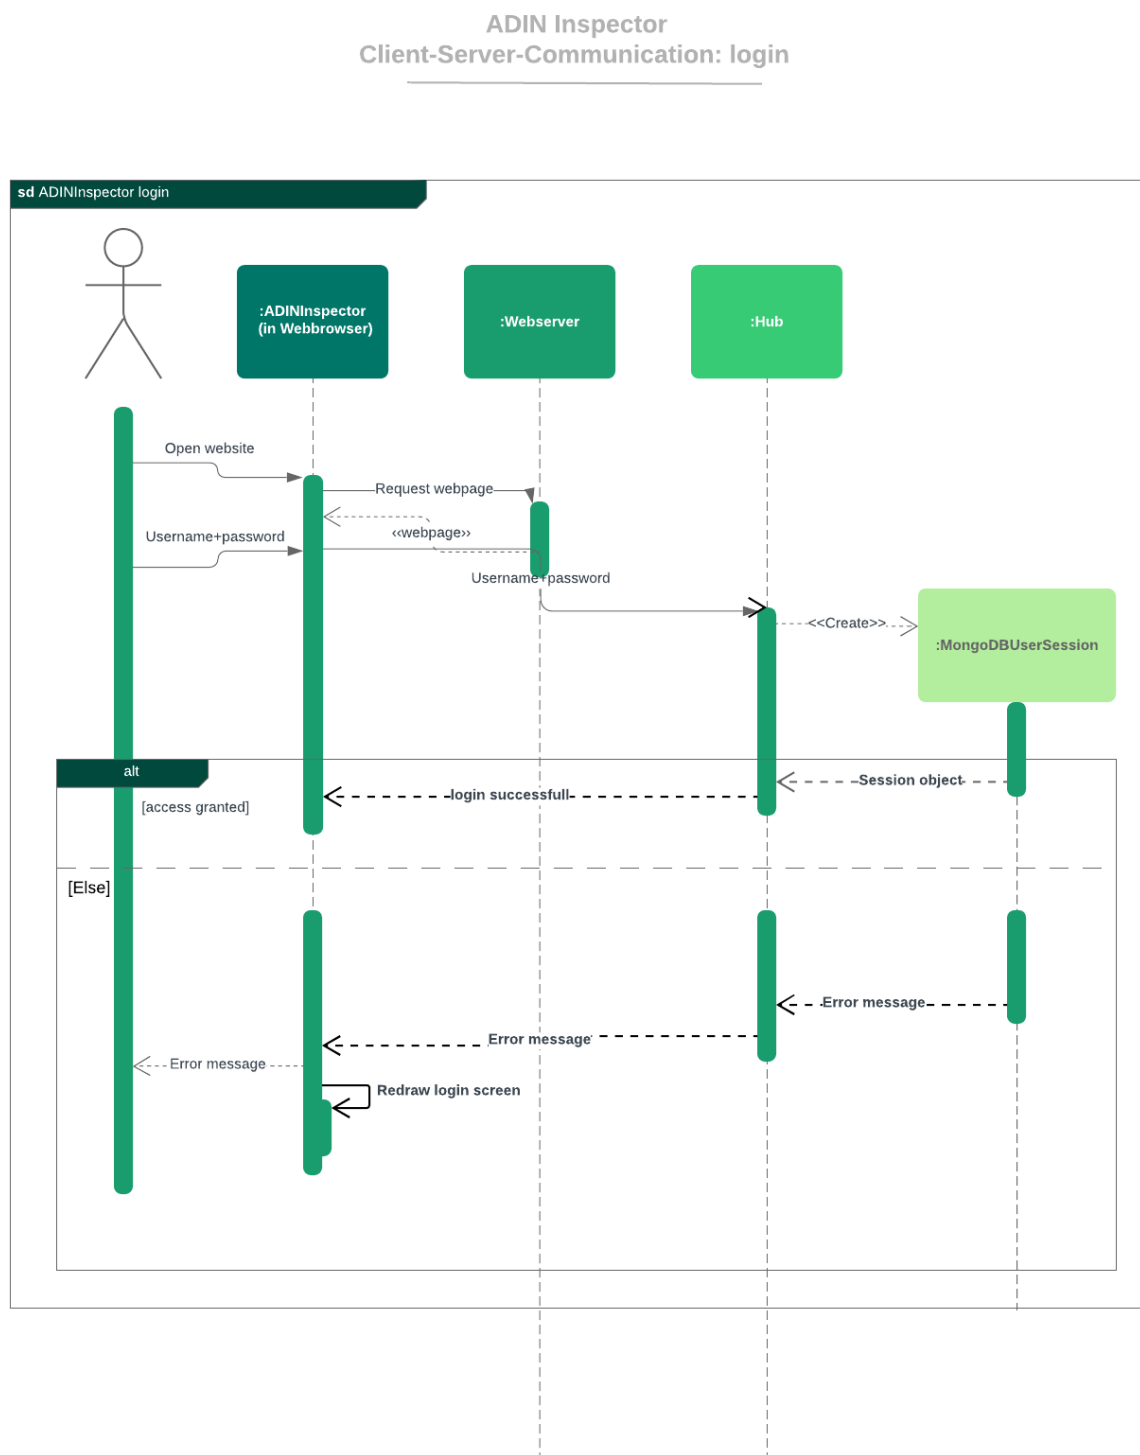
1 Design

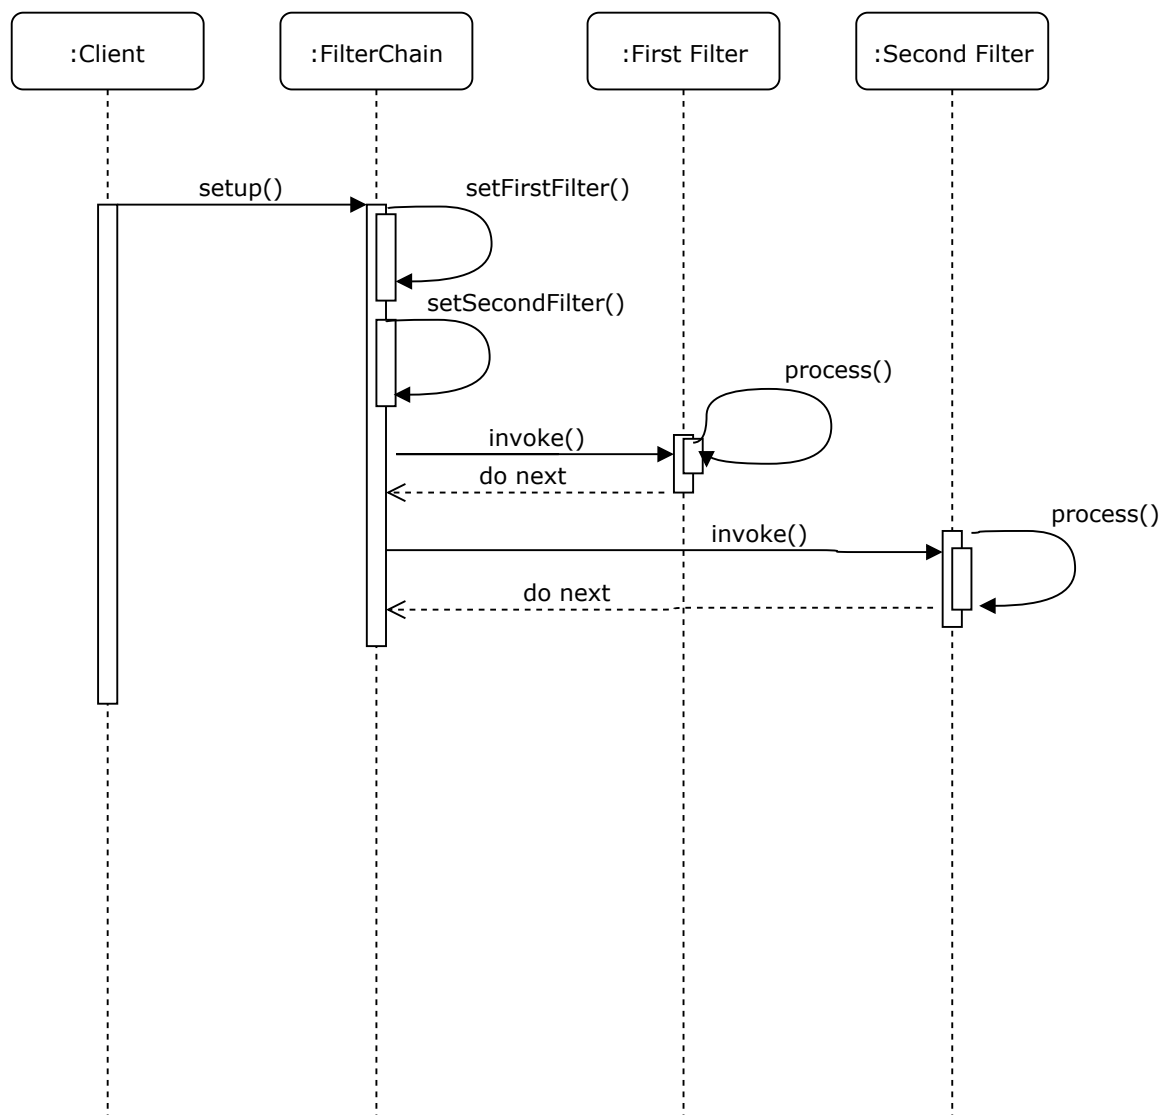
1.1 Front-End

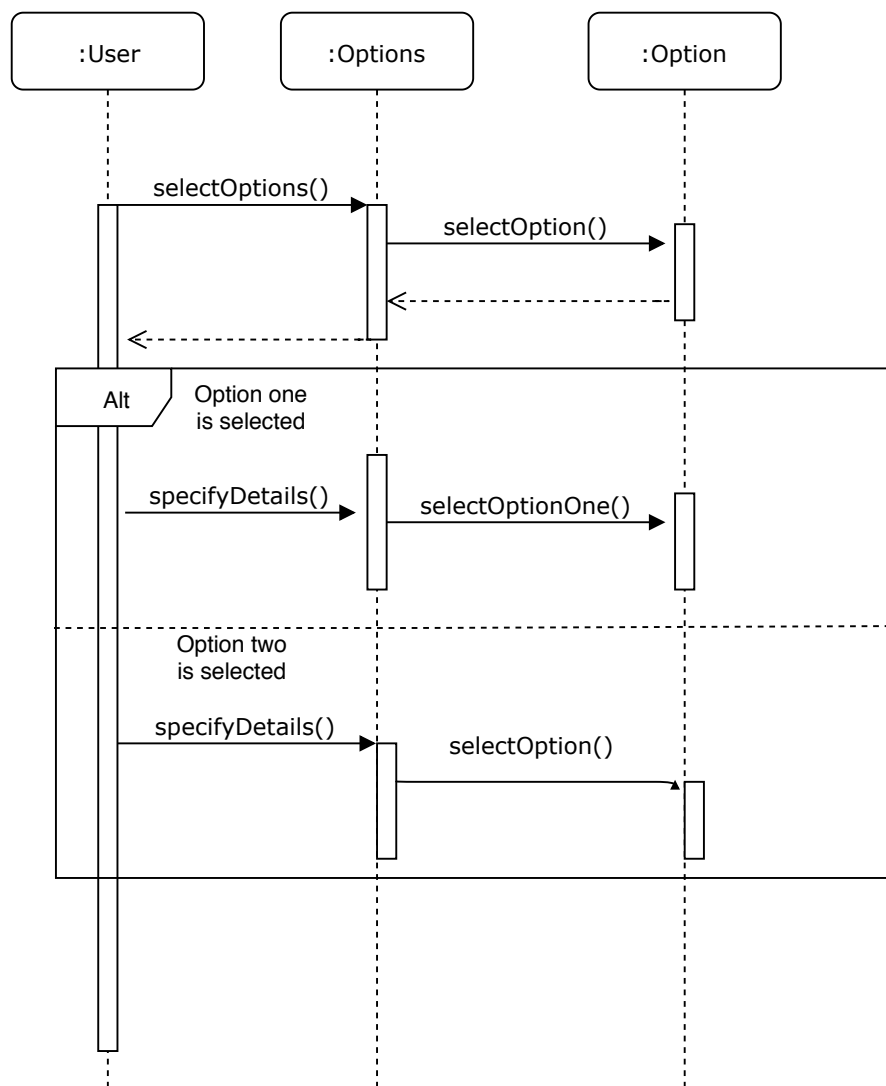
1.1.1 Sequence Diagram



The following diagram shows an alternative view of the login sequence:







This diagram shows the control flow for handling a movement of the slider by the user.

1.1.2 Activity Diagram

1.1.3 Class Diagram

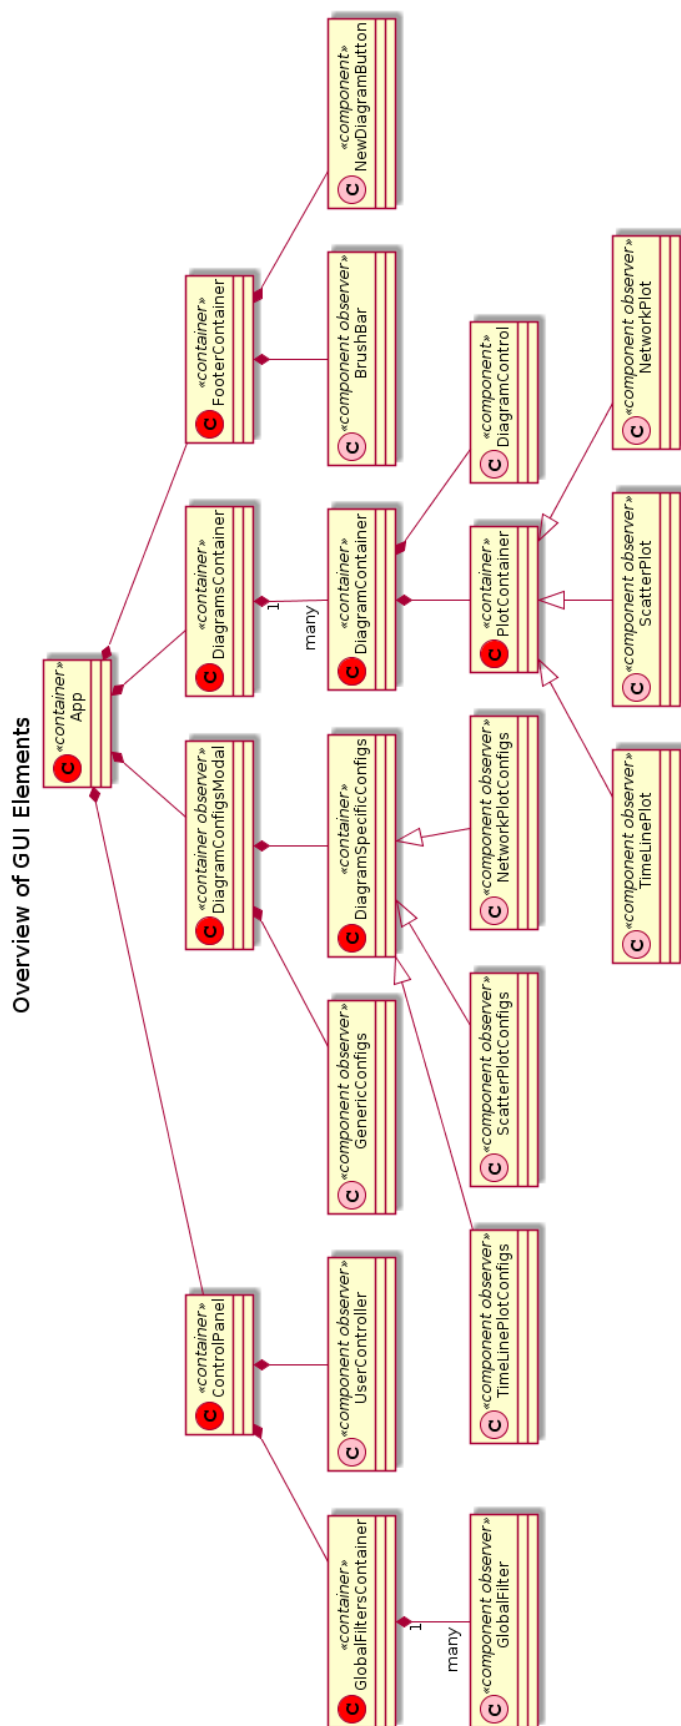


Figure 1: This diagram shows an overview of GUI elements of the main application, when the user has successfully logged in.

1.2 Client-server protocol

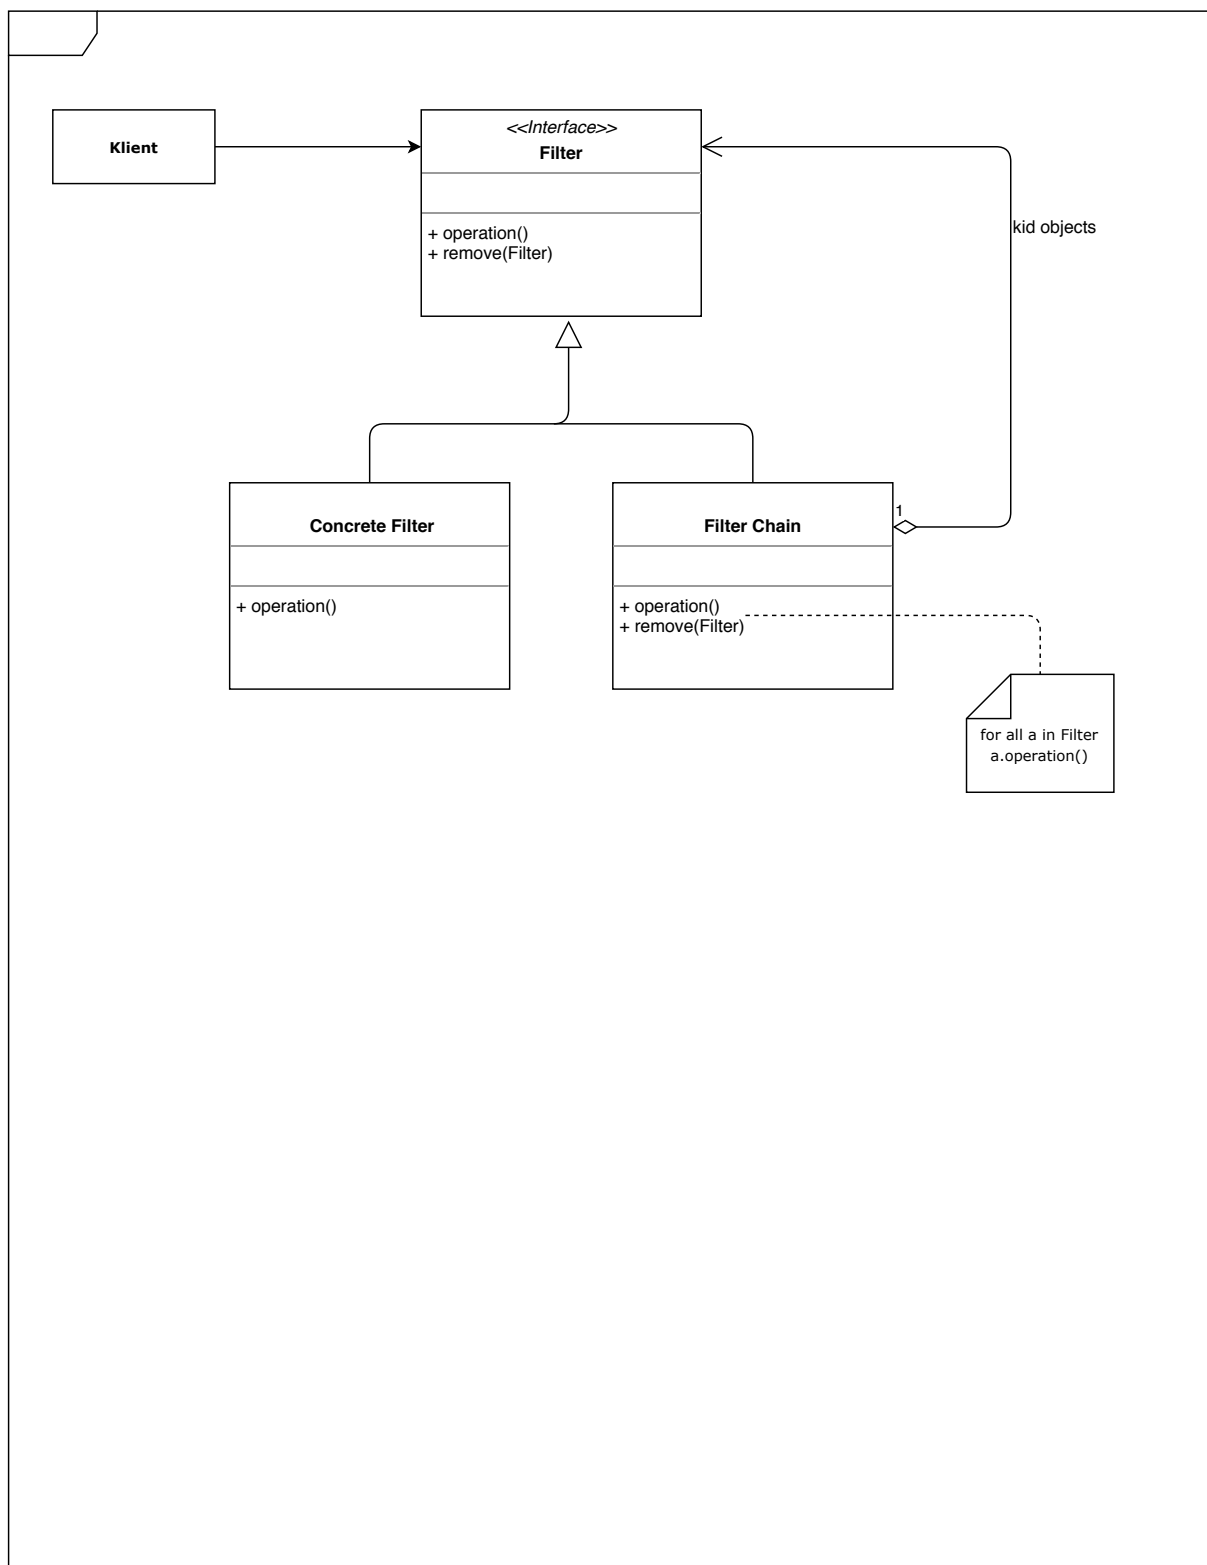
Messages between client and server are exchanged as strings in JSON format. In the following list words in angle brackets ("**<>**") are placeholders.

1.2.1 Requests from client to server:

- `getAvailableCollections`
syntax: {"cmd": "GET_AV_COLL"}
expected response: list of collections
- `getCollectionSize(collection)`
syntax: {"cmd": "GET_COLL_SIZE", "par": "<collection>"}
where <collection> is the name of a collection
expected response: collection size
- `getCollection(collection)`
syntax: {"cmd": "GET_COLL", "par": "<collection>"}
expected response: data set
- `getRecordsInRange(collection, key, start, end)`
syntax: {"cmd": "GET_RECORDS_RANGE", "par": "<collection>", "key": "<keyvalue>", "start": "<startvalue>", "end": "<endvalue>"}
where <key> is the name of a key in the given collection and <startvalue> and <endvalue> are valid values for this key
expected response: data set
- `getRecordsInRangeSize(collection, key, start, end)`
syntax: {"cmd": "GET_RECORDS_RANGE_SIZE", "par": "<collection>", "key": "<key-value>", "start": "<startvalue>", "end": "<endvalue>"}
expected response: collection size

1.2.2 Messages from server to client:

- list of collections
syntax: {"cmd": "LIST_COL", "par": ["<collection>"]}
where <collection> is the name of a collection
- collection size
syntax: {"cmd": "COLL_SIZE", "par": "<size>"}
where <size> is the number of records in this collection
- data set
syntax: {"cmd": "DATA", "par": [<record>]}
where each record is a JSON object



1.3 Back-End

This subsection deals with the back-end of the ADIN INSPECTOR. How the system deals with client http calls, and how kafka interacts with the system. An overview of the system can be seen in Figure 2

1.3.1 Class Diagram

Next we'll look at each class and method in detail

- **Config properties file**
The config file is stored alongside the built application .jar file and contains the path to the Kafka installation folder, the user name and password of a mongoDB account with the highest level of access and the name of the database.
- **Initializer**
Methods:
 - main
parameters: String of arguments from the console
returns: void
App entry point.
We load the config.properties file and use the path provided to start the zookeeper, kafka and mongodb services
- **MongoConsumer**
The Mongo Consumer, as the name implies, consumes all messages from all topics in the Kafka messaging system. Once a message is found it is passed along to the Mongo Client for further processing.
Variables
 - clientMediator
Type : MongoClientMediator
An instance of the Mongo Client Mediator, created with the credentials from the config file.
 Methods
 - MongoConsumer constructor
parameters: user name and password of a mongoDB account with the highest level of access.
Initializes the MongoClient variable and calls listenForRecords();
 - getAllTopics
parameters: none
returns: an array of strings containing all the available kafka Topics.
Asks the kafka server service which topics exists.
 - listenForRecords
parameters: none

returns: void

This Method first calls getAllTopics and uses the array of topics to poll the kafka server for new messages.

If new messages are found then the messages are passed to the Mongo Mediator for adding them to the Database.

If no new messages are found for a topic notify the Mongo Mediator that the collection tied to the topic is ready for pre-processing.

- MongoClientMediator This object serves as a nexus between the users who want to get data out of the database and the consumer and dataProcessor who want to add data into the database. This class encapsulates the mongo client from the mongo API.

Variables

- client

type: MongoClient

An instance of the Mongo Client from the official java API.

- dataProc

A reference to the data processor class for this client.

Methods

- MongoClientMediator constructor

parameters: Username and password

Initializes the client variable, throws an error if the user is not found.

- addRecordToCollection

parameters: String representation of a record in json format

String name of the collection it should be added to.

returns: void

Converts the json string into a java object, then to a bson document and uses the mongoAPI to insert it into the database.

- addRecordsToCollection

parameters: String Array of records to be added to a collection

String name of the collection it should be added to.

returns: void

for each one of the members of the array call addRecordToCollection

- ProcessCollection

parameters: String, name of a collection

returns: void

signal the data processor to start the processing of a collection

- getCollection

parameters: String, name of a collection

returns: String array containing all entries of the collection

- getStartRecord

parameters: String, name of a collection

returns: the first entry of the collection as a String.

- getEndRecord
parameters: String, name of a collection
returns: the last entry of the collection as a String.
 - getCollectionSize
parameters: String, name of a collection
returns: the number of entries in the collectoin as int
 - getRecordsInRange
parameters: String, name of the collection to query
String, key of the parameter used for filtering
String start and end ranges for the filtering
returns: String array containing all entries of the collection within that range
this Method is very general to allow for flexibility. For example by letting the key be, SourceIPAddresses, or a timeStamp.
 - getRecordsInRangeSize
parameters: String, name of the collection to query
String, key of the parameter used for filtering
String start and end ranges for the filtering
returns: number of elements matching the range as int
- Record
Every message that comes from kafka and needs to be added to the database has it's own Record class that inherit from this one.
Every single class that inherits needs to be able to, using reflection, convert itself into a Bson Document where every variable is a key Value pair of the name of the variable and it's associated value.
Variables
 - id
type: String
 Methods
 - getAsDocument()
parameters: none
returns: A Document, containing every variable of any class inheriting from this one.
This function checks for every variable, gets it's name and value as a string and adds it to the document that it eventually returns.
 - PacketRecord
Inheriting from Record, this class contains the variables that match the json string obtained from kafka.
Variables
 - id
type: String

this id is used for determining the ordering when saving to mongoDB, it's the offset of the message in the kafka messaging queue. inherited from Record

- client
type: String
- L2Protocol
type: String
- SourceMACAddress
type: String
- L4Protocol
type: String
- SourceIPAddress
type: String
- PacketSummary
type: String
- DestinationIPAddress
type: String
- Timestamp
type: String
- DestinationPort
type: String
- SourcePort
type: String
- DestinationMACAddress
type: String

Methods

- getters / setters
parameters: variable
returns: variable type
Each variable has it's getters and setter methods.

- AlarmRecord

Inheriting from Record, this class contains the variables that match the json string obtained from kafka.

Variables

- id
type: String
- AlarmID
type: String
- AlarmType
type: String

- AlarmOccurrenceTime
type: String
- AlarmCategory
type: String
- AlarmScore
type: String
- AlarmDescription
type: String
- PacketSummary
type: String

Methods

- getters / setters
parameters: variable
returns: variable type
Each variable has it's getters and setter methods.

- MiscRecord

Inheriting from Record, this class is used by the data processor as an 'in-between' state before saving to the database. As well as an extension point for adding more types of records into the database programatically in the future.

Refer to the data processor class for further data on the key value pairs.

Variables

- pairs
A Map of strings to Objects to store any 1 to many relationships

Methods

- getters / setters
parameters: none
returns: variable type
Each variable has it's getters and setter methods.

- DataProcessor

This class is a mediator for each one of our data aggregators used for extraciton of features from the raw data stored in mongoDB.

We might want to hve multiple data processors for chaining different aggregators together or to split up the work into mutliple threads. This is dependant on further performance testing.

Variables

- client
an instance of the associated mongoClient that requested the data aggregation
- aggregators
A Arraylist containing all the aggregators to be applied on a collection.

Methods

- getters / setters
- processData
 - parameters: variable
 - returns: variable type

• IAggregator

This interface is the building block for every aggregator to be applied to data

Variables Methods

- processData
 - parameters: Records array of the records to be processed

• FlowRatePerSecond

Implements IAggregator. This calculates, per port, the outgoing and ingoing connections. A record processed by this aggregator is stored in a collection as follows:

Name of collection: `collectionName_FlowratePerSec`

structure of record as json:

```
{
  "date" : \{" date" " Unix_Timestamp  }
  rounded down to the second this record points to.
  Connections : [
    { Port: "portNumber", "InOut" : " In/Out ", count : "Number" }
    { Port: "portNumber", "InOut" : " In/Out ", count : "Number" }
    ...
  ]
}
```

] This array has an entry per port if the port communicated that second.

Precomputing this allows us to stream whenever the client needs the information for a specific node.

}

Methods

- processData
 - parameters: Records array of the records to be processed
 - specific implementation left to the classes implementing this interface

• NumberOfConnectionsPerNodePerSecond

Implements IAggregator. This calculates the outgoing and ingoing connections. A record processed by this aggregator is stored in a collection as follows:

Name of collection: `collectionName_FlowratePerSec`

structure of record as json:

```
{
  "date" : \{" date" " Unix_Timestamp  }
  rounded down to the second this record points to.
```



```

Connections : [
{ Port: "portNumber", count : "Number" }
{ Port: "portNumber", count : "Number" }
...

```

] This array has an entry per port if the port communicated that second.
 Precomputing this allows us to stream whenever the client needs the information for a specific node.

```

}
```

Methods

- processData
 parameters: Records array of the records to be processed

• Hub

This class implements the network handlers for the websocket connection to the client and access methods for a database connection.

Variables

- requestHandler
 Type : ClientProtocolHandler
 The strategy object we call for the actual parsing of the client requests.
- database
 Type : IUserSession
 The database we use during a user session.

Methods

- handleOpen
 parameters: Session session - the current session
 returns: void
 Event handler for the start of websocket connection.
- handleClose
 parameters: Session session - the current session
 returns: void
 Event handler for closing a connection.
- handleMessage
 parameters: String message - the message that we received from the client
 Session session - the current session
 returns: String - the response to be sent to the client
 Event handler for receiving a message. The message is passed to the ClientProtocol-Handler.
- handleError
 parameters: Session session - the current session
 Throwable t - the exception that occurred

returns: void

Event handler for errors/exceptions during communication.

- IUserSession

An IUserSession object encapsulates a data base session. On instantiation an IUserSession connects to a database using the given user id and password and uses this connection for all following data base access.

Methods

- UserSession

parameters: String username - the user id to login with

String password - the password

returns: IUserSession

Factory method to instantiate a new UserSession and log in into the database using the given credentials.

- getAvailableCollections

parameters: -

returns: String array with collection names

Returns an array with the names of the collections available to the current user.

- getCollectionSize

parameters: String collection - the collection to query

returns: long - the number of records

Returns the number of records in the specified collection.

- getCollection

parameters: String - name of a collection

returns: String array containing all entries of the collection

- getRecordsInRange

parameters: String - name of the collection to query

String key - the parameter used for filtering

String start and end - range for the filtering

returns: String array containing all entries of the collection within the filter range

Returns an array containing all records of this collection for which the value of the specified key is in the range [start, end). The records will be in the same order as they are in the collection.

- getRecordsInRangeSize

parameters: String - name of the collection to query

String key - the parameter used for filtering

String start and end - range for the filtering

returns: number of elements matching the range as int

Returns the number of records in the specified collection for which the value of the specified key is within the range [start, end).

- MongoDBUserSession

Encapsulates a user session for a connection to a MongoDB database.

Methods

- MongoDBUserSession constructor
parameters: -
Private constructor to create a new MongoDB session.
- UserSession
parameters: String username - the user id to login with
String password - the password
returns: a new MongoDBUserSession object
Factory method to instantiate a new MongoDBUserSession and log in into the database using the given credentials.
- getAvailableCollections
parameters: -
returns: String array with collection names
Returns an array with the names of the collections available to the current user.
- getCollectionSize
parameters: String collection - the collection to query
returns: long - the number of records
Returns the number of records in the specified collection.
- getCollection
parameters: String - name of a collection
returns: String array containing all entries of the collection
- getRecordsInRange
parameters: String - name of the collection to query
String key - the parameter used for filtering
String start and end - range for the filtering
returns: String array containing all entries of the collection within the filter range
Returns an array containing all records of this collection for which the value of the specified key is in the range [start, end). The records will be in the same order as they are in the collection.
- getRecordsInRangeSize
parameters: String - name of the collection to query
String key - the parameter used for filtering
String start and end - range for the filtering
returns: number of elements matching the range as int
Returns the number of records in the specified collection for which the value of the specified key is within the range [start, end).

- ClientRequestHandler

This class handles client requests by parsing them, executing the requested action and producing responses. The requested actions are typically executed by calls to the database session object.

Methods

- handleRequest

parameters:

IUserSession dbSession - the current database session


Session session - the current client session

String message - the client request to process

returns: String - the response to be sent to the client

Parse the message from the client, execute the requested action (typically a database query) and construct the response message.

1.3.2 Sequence Diagram



1.3.3 Activity Diagram